

SSC 335/394: Scientific and Technical Computing

Computer Architectures: parallel computers



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

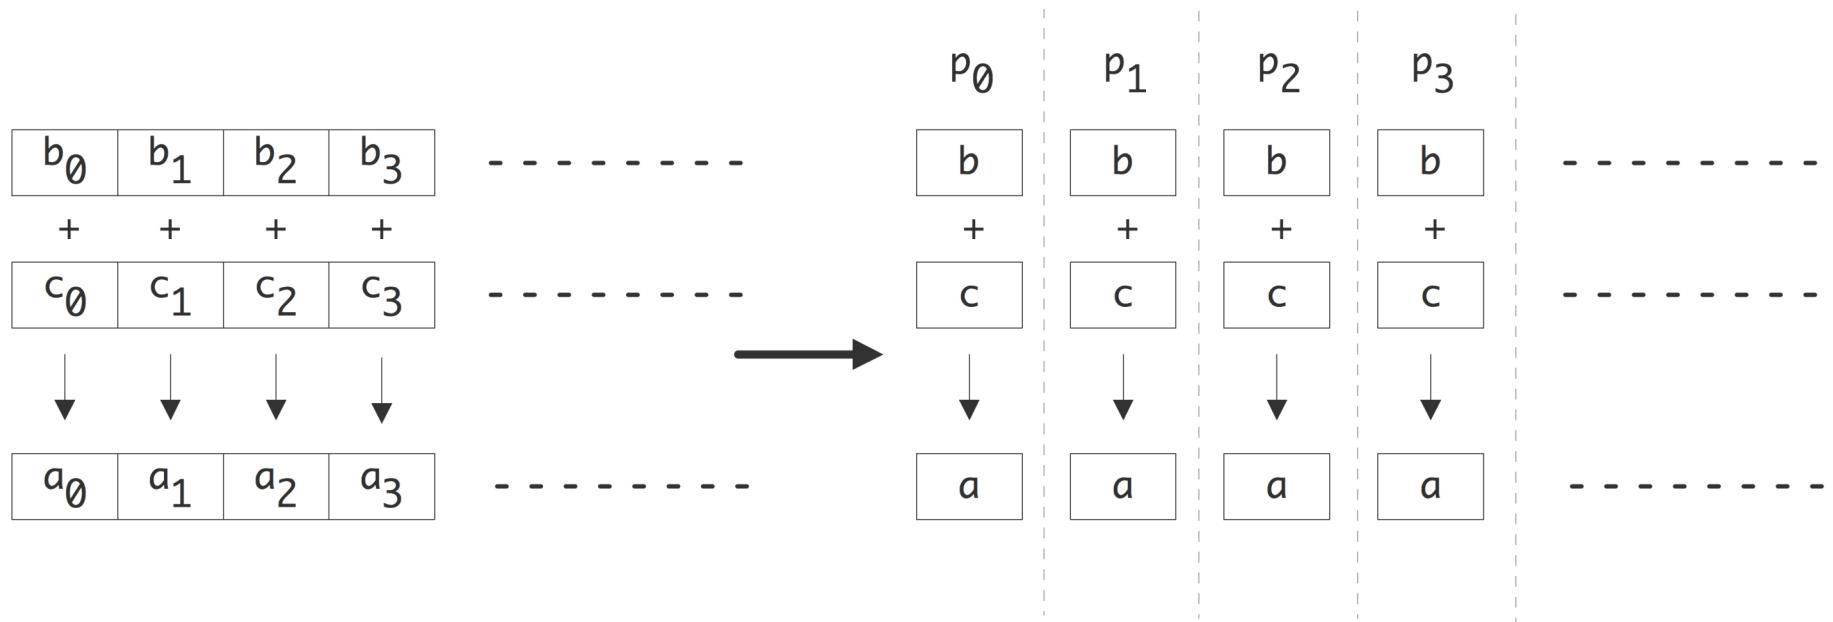
The basic idea

- Spread operations over many processors
- If n operations take time t on 1 processor,
- Does this become t/p on p processors ($p \leq n$)?

```
for (i=0; i<n; i++)
  a[i] = b[i]+c[i]
```

```
a = b+c
```

Idealized version: every process has one array element



The basic idea

- Spread operations over many processors
- If n operations take time t on 1 processor,
- Does this become t/p on p processors ($p \leq n$)?

```
for (i=0; i<n; i++)
  a[i] = b[i]+c[i]
```

```
a = b+c
```

```
for (i=my_low; i<my_high; i++)
  a[i] = b[i]+c[i]
```

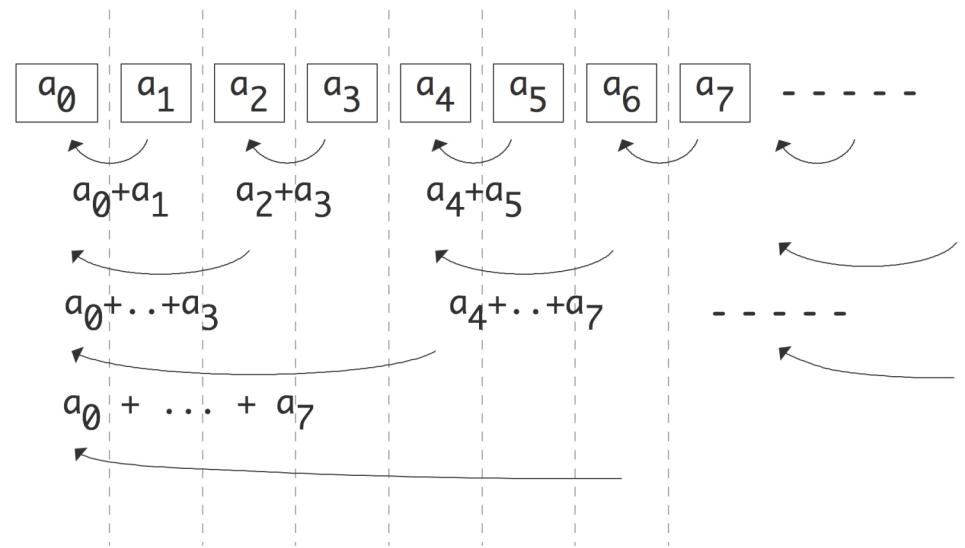
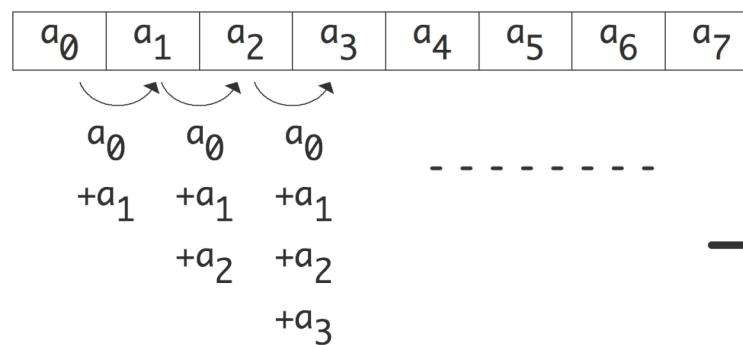
Idealized version: every process has one array element

Slightly less ideal: each processor has part of the array

The basic idea (cont'd)

- Spread operations over many processors
- If n operations take time t on 1 processor,
- Does it always become t/p on p processors ($p \leq n$)?

```
s = sum( x[i], i=0,n-1 )
```



The basic idea (cont'd)

- Spread operations over many processors
- If n operations take time t on 1 processor,
- Does it always become t/p on p processors ($p \leq n$)?

```
s = sum( x[i], i=0, n-1 )
```

```
for (p=0; p<n/2; p++)
    x[2p,0] = x[2p]+x[2p+1]
for (p=0; p<n/4; p++)
    x[4p,1] = x[4p]+x[4p+2]
for ( .. p<n/8 .. )
```

Et cetera

Conclusion: n operations can be done with $n/2$ processors, in total time $\log_2 n$

Theoretical question: can addition be done faster?

Practical question: can we even do this?

Some theory

-before we get into the hardware
- Optimally, P processes give $T_P = T_1/P$
- Speedup $S_P = T_1/T_p$, is P at best
- Superlinear speedup not possible in theory, sometimes happens in practice.
- Perfect speedup in “embarrassingly parallel applications”
- Less than optimal: overhead, sequential parts, dependencies

Some more theory

-before we get into the hardware
- Optimally, P processes give $T_p = T_1/P$
- Speedup $S_p = T_1/T_p$, is P at best
- Efficiency $E_p = S_p/P$
- Scalability: efficiency bounded below

Scaling

- Increasing the number of processors for a given problem makes sense up to a point: $p > n/2$ in the addition example has no use
- Strong scaling: problem constant, number of processors increasing
- More realistic: scaling up problem and processors simultaneously, for instance to keep data per processor constant: Weak scaling
- Weak scaling not always possible: problem size depends on measurements or other external factors.

Amdahl's Law

- Some parts of a code are not parallelizable
- => they ultimately become a bottleneck
- For instance, if 5% is sequential, you can not get a speedup over 20, no matter P.
- Formally: $F_p + F_s = 1$, $T_p = T_1(F_s + F_p/p)$, so T_p approaches $T_1 F_s$ as p increases

More theory of parallelism

- PRAM: Parallel Random Access Machine
- Theoretical model
 - Not much relevance to practice
 - Often uses (implicitly) unrealistic machine models

Theoretical characterization of architectures



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

Parallel Computers Architectures

- **Parallel computing** means using multiple processors, possibly comprising multiple computers
- Flynn's (1966) taxonomy is a first way to classify parallel computers into one of four types:
 - (SISD) Single instruction, single data
 - Your desktop (unless you have a newer multiprocessor one)
 - (SIMD) Single instruction, multiple data:
 - Thinking machines CM-2
 - Cray 1, and other vector machines (there's some controversy here)
 - Parts of modern GPUs
 - (MISD) Multiple instruction, single data
 - Special purpose machines
 - No commercial, general purpose machines
 - (MIMD) Multiple instruction, multiple data
 - Nearly all of today's parallel machines

SIMD

- Based on regularity of computation: all processors often doing the same operation: *data parallel*
- Big advantage: processor do not need separate ALU
- ==> lots of small processors packed together
- Ex: Goodyear MPP: 64k processors in 1983
- Use masks to let processors differentiate

SIMD then and now

- There used to be computers that were entirely SIMD (usually attached processor to a front end)
- SIMD these days:
 - SSE instructions in regular CPUs
 - GPUs are SIMD units (sort of)

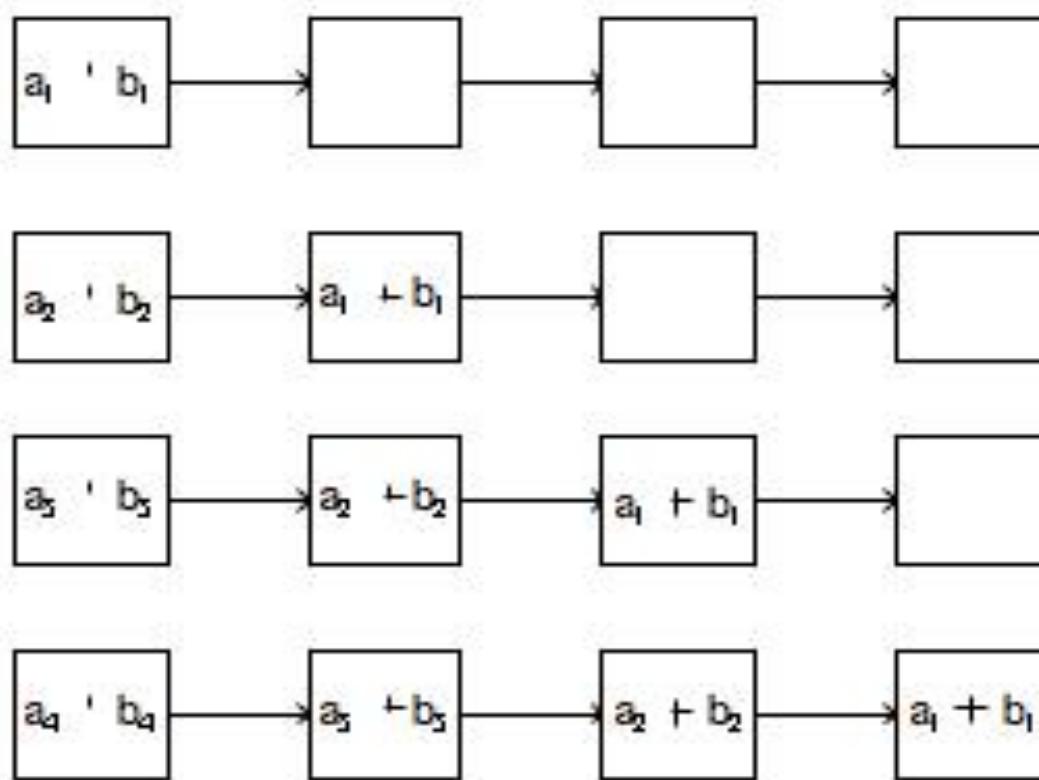
Kinda SIMD: Vector Machines

- Based on a single processor with:
 - Segmented (pipeline) functional units
 - Needs sequence of the same operation
- Dominated early parallel market
 - overtaken in the 90s by clusters, et al.
- Making a comeback (sort of)
 - clusters/constellations of vector machines:
 - Earth Simulator (NEC SX6) and Cray X1/X1E
 - Arithmetic units in CPUs are pipelined.

Pipeline

- Assembly line model (body on frame, attach wheels, doors, handles on doors)
- Floating point multiply: exponent align,multiply, exponent normalize
- Separate hardware for each stage: pipeline processor

$$c_1 \leftarrow a_1 + b_1$$



Pipeline'

- Complexity model: asymptotic rate, $n^{1/2}$
- Multi-vectors, parallel pipes (demands on code)
- Is like SIMD
- (There is also something called an “instruction pipeline”)
- Requires independent operations:
 $a_i \leq b_i + c_i$
not:
 $a_i \leq b_i + a_{i-1}$

MIMD

- Multiple Instruction, Multiple Data
- Most general model: each processor works on its own data with its own data stream: *task parallel*
- Example: one processor produces data, next processor consumes/analyzes data

MIMD

- In practice SPMD: Single Program Multiple Data:
 - all processors execute the same code
 - Just not the same instruction at the same time
 - Different control flow possible too
 - Different amounts of data: load unbalance

Granularity

- You saw data parallel and task parallel
- Medium grain parallelism: carve up large job into tasks of data parallel work
- (Example: array summing, each processor has a subarray)
- Good match to hybrid architectures:
task -> node
data parallel -> SIMD engine

GPU: the miracle architecture

- Lots of hype about incredible speedup / high performance for low cost. What's behind it?
- Origin of GPUs: that “G”
- Graphics processing: identical (fairly simple) operations on lots of pixels
- Doesn't matter when any individual pixel gets processed, as long as they all get done in the end
- (Otoh, CPU: heterogeneous instructions, need to be done ASAP.)
- => GPU is SIMD engine
- ...and scientific computing is often very data-parallel

GPU programming:

- KernelProc<< m,n >>(args)
- Explicit SIMD programming
- There is more: threads (see later)

Characterization by Memory structure

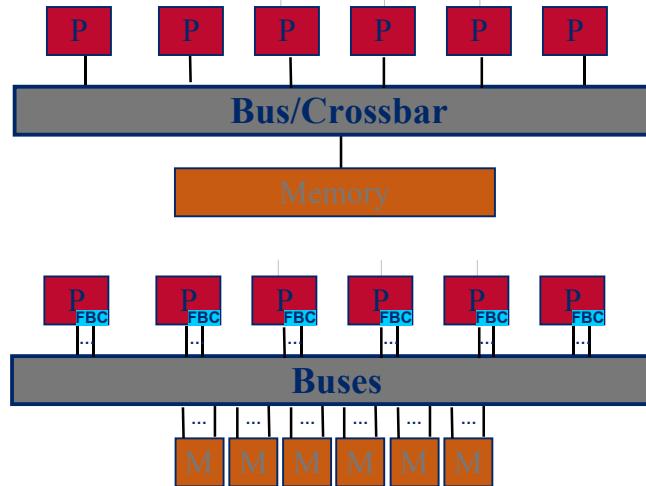
Parallel Computer Architectures

- Top500 List now dominated by MPPs and Clusters
- The MIMD model “won”.
- SIMD exists only on smaller scale
- A much more useful way to classification is by memory model
 - *shared* memory
 - *distributed* memory

Two memory models

- Shared memory: all processors share the same address space
 - OpenMP: directives-based programming
 - PGAS languages (UPC, Titanium, X10)
- Distributed memory: every processor has its own address space
 - MPI: Message Passing Interface

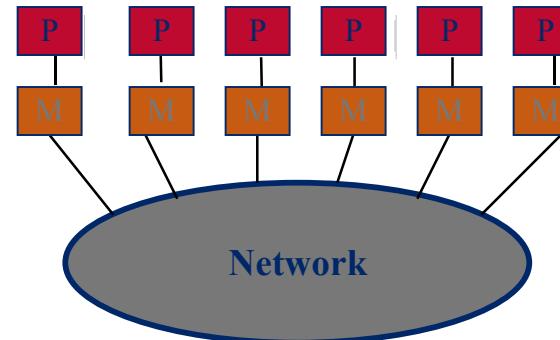
Shared and Distributed Memory



Shared memory: single address space. All processors have access to a pool of shared memory.
(e.g., Single Cluster node (2-way, 4-way, ...))

Methods of memory access :

- Bus
- Distributed Switch
- Crossbar

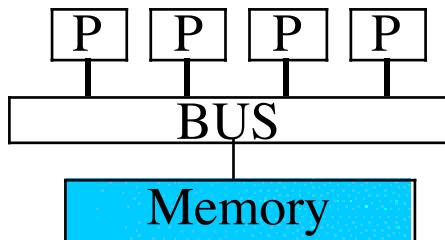


Distributed memory: each processor has its own local memory. Must do message passing to exchange data between processors.
(examples: Linux Clusters, Cray XT3)

Methods of memory access :

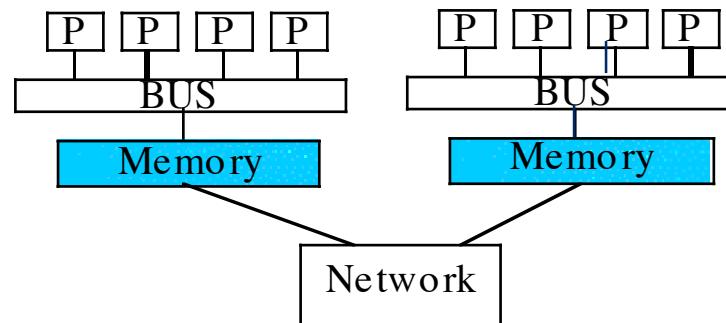
- single switch or switch hierarchy with fat tree, etc. topology

Shared Memory: UMA and NUMA

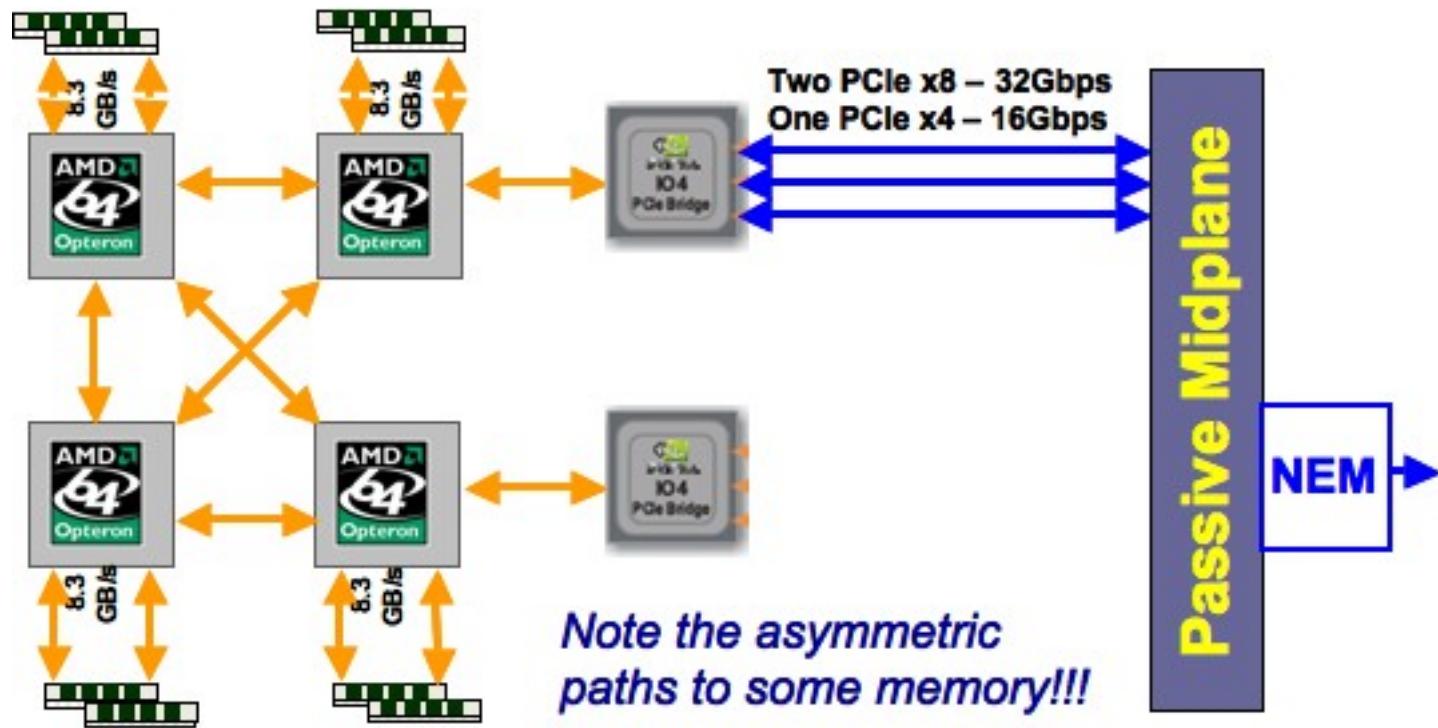


Uniform Memory Access (UMA):
Each processor has uniform access time to memory - also known as symmetric multiprocessors (SMPs)
(example: Sun E25000 at TACC)

Non-Uniform Memory Access (NUMA):
Time for memory access depends on location of data; also known as Distributed Shared memory machines.
Local access is faster than non-local access. Easier to scale than SMPs
(e.g.: SGI Origin 2000)



Numa example: ranger node



Interconnects



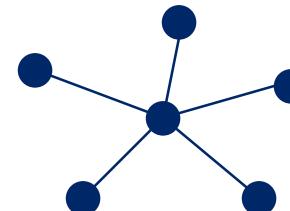
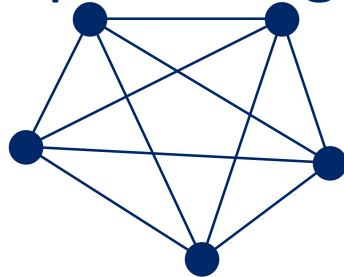
THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

Topology of interconnects

- What is the actual ‘shape’ of the interconnect? Are the nodes connect by a 2D mesh? A ring? Something more elaborate?
- => some graph theory

Completely Connected and Star Networks

- Completely Connected : Each processor has direct communication link to every other processor (compare ranger node)



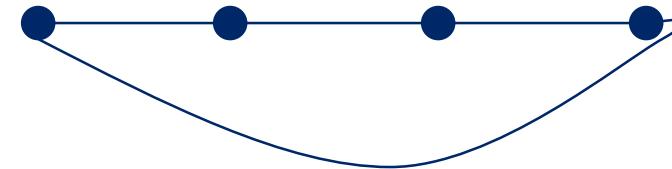
- Star Connected Network : The middle processor is the central processor; every other processor is connected to it.

Arrays and Rings

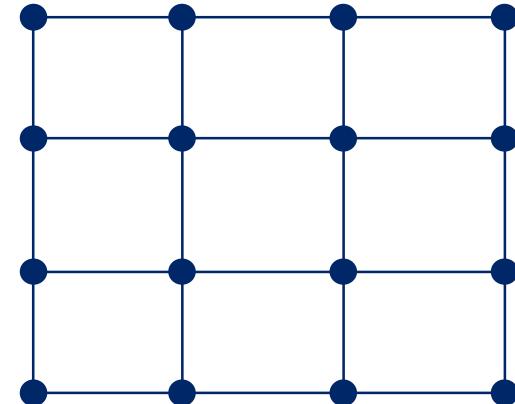
- Linear Array :



- Ring :

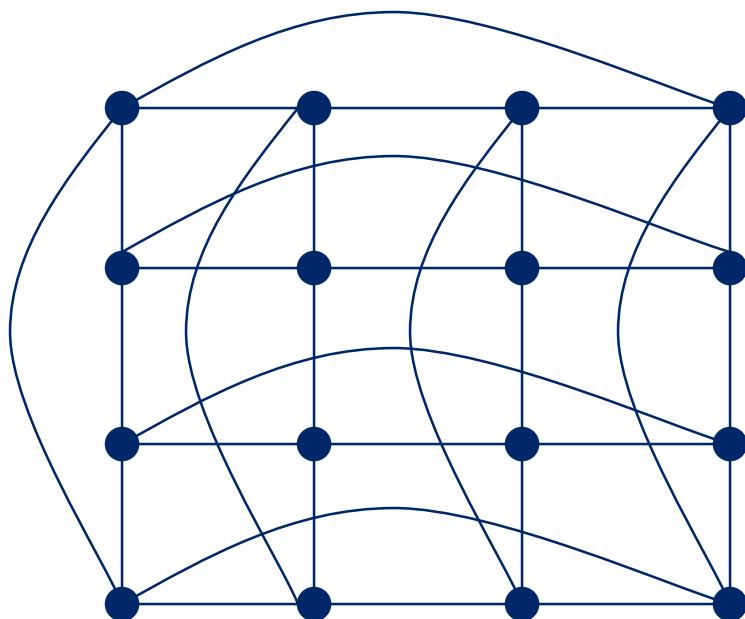


- Mesh Network (e.g. 2D-array)



Torus

2-d Torus (2-d version of the ring)

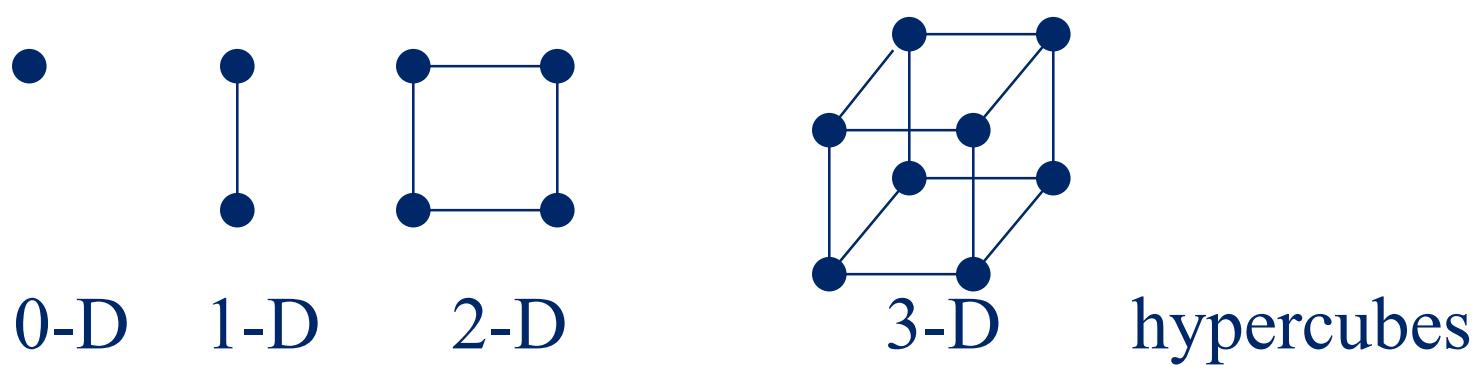


Hypercubes

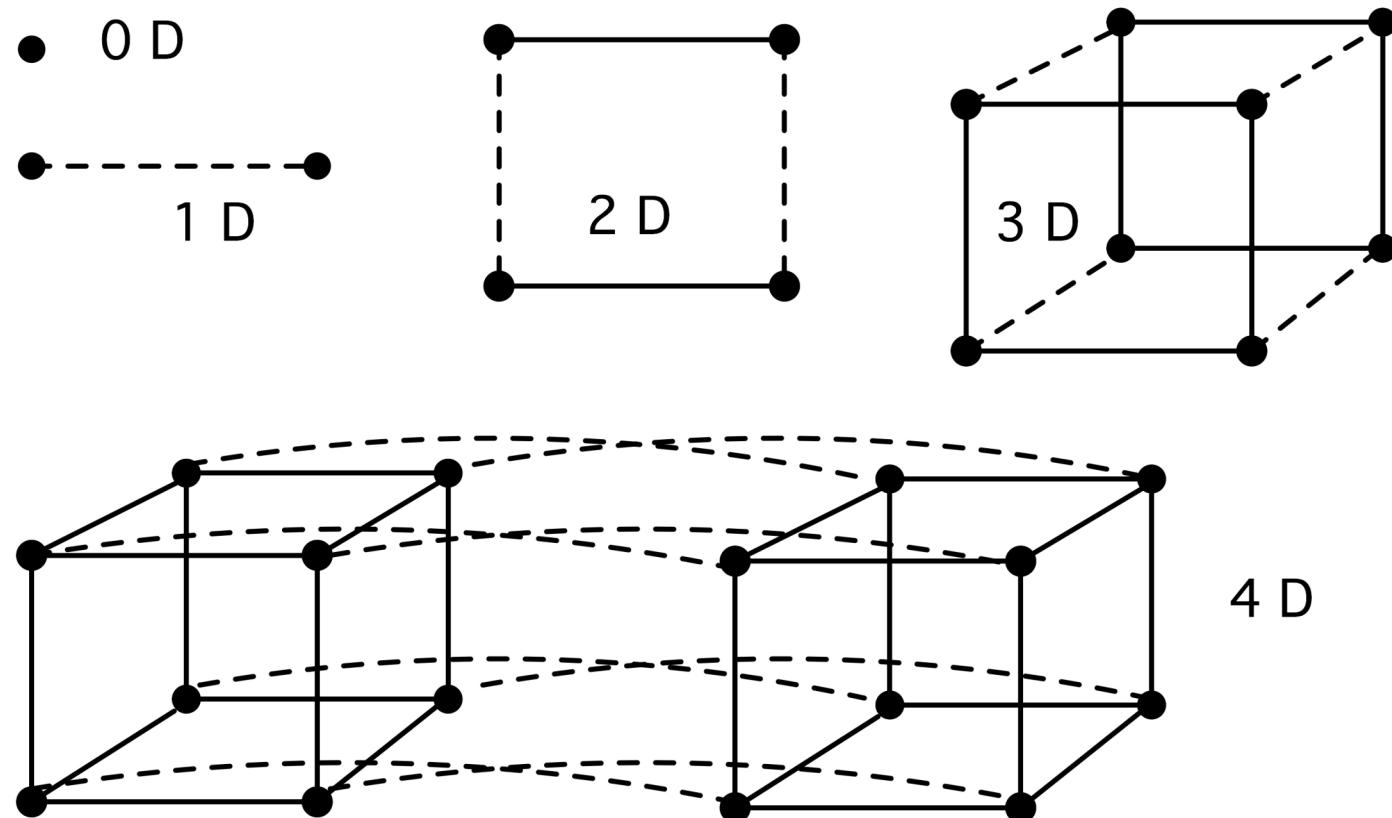
- Hypercube Network : A multidimensional mesh of processors with exactly two processors in each dimension. A d dimensional processor consists of

$$p = 2^d \text{ processors}$$

- Shown below are 0, 1, 2, and 3D hypercubes



Inductive definition

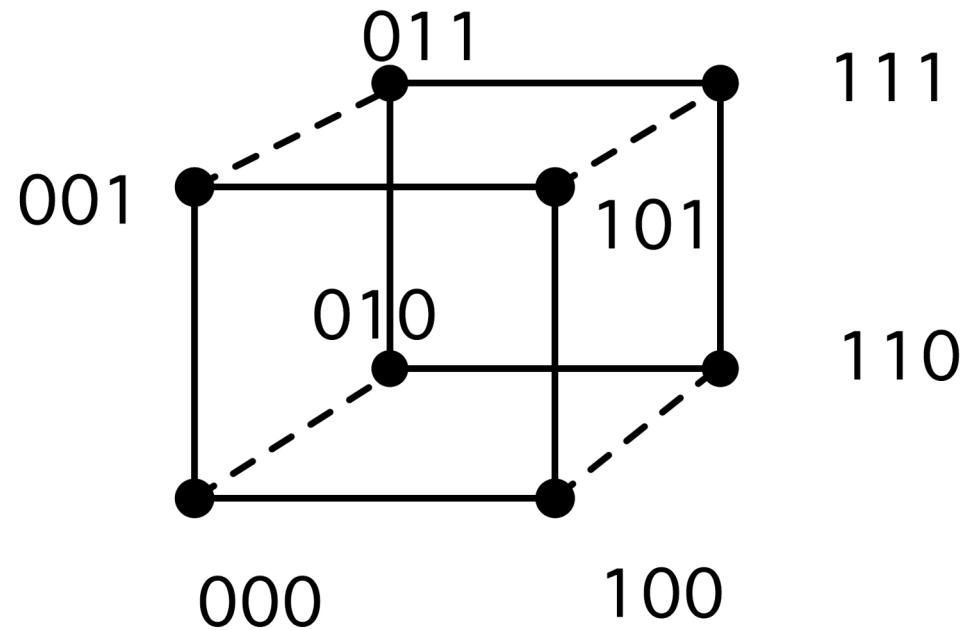


Pros and cons of hypercubes

- Pro: processors are close together: never more than $\log(P)$
- Lots of bandwidth
- Little chance of contention
- Con: the number of wires out of a processor depends on P: complicated design
- Values of P other than 2^p not possible.

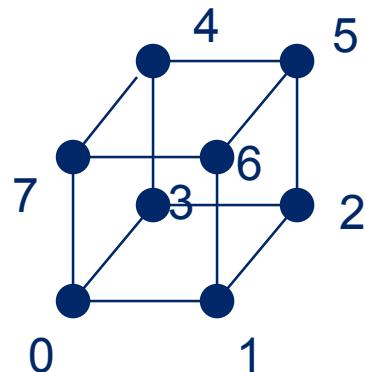
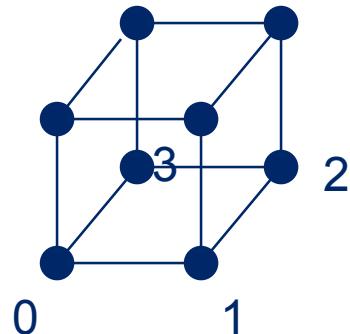
Mapping applications to hypercubes

- Is there a natural mapping from 1,2,3D to a hypercube?
- Naïve node numbering does not work:
- Nodes 0 and 1 have distance 1, but
- 3 and 4 have distance 3
- (so do 7 and 0)



Mapping applications to hypercubes

- Is there a natural mapping from 1,2,3D to a hypercube?
- => Gray codes
- Recursive definition: number subcube, then other subcube in mirroring order.



Subsequent processors (in the Linear ordering) all one link apart

Recursive definition:

0 | 1

0 0 | 1 1

0 1 | 1 0

0 0 0 | 1 1 1 1

0 0 1 1 | 1 1 0 0

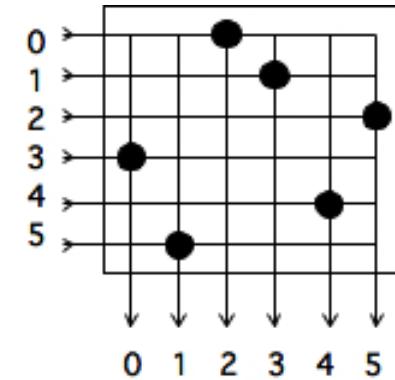
0 1 1 0 | 0 1 1 0

Busses/Hubs and Crossbars

Hub/Bus: Every processor shares the communication links

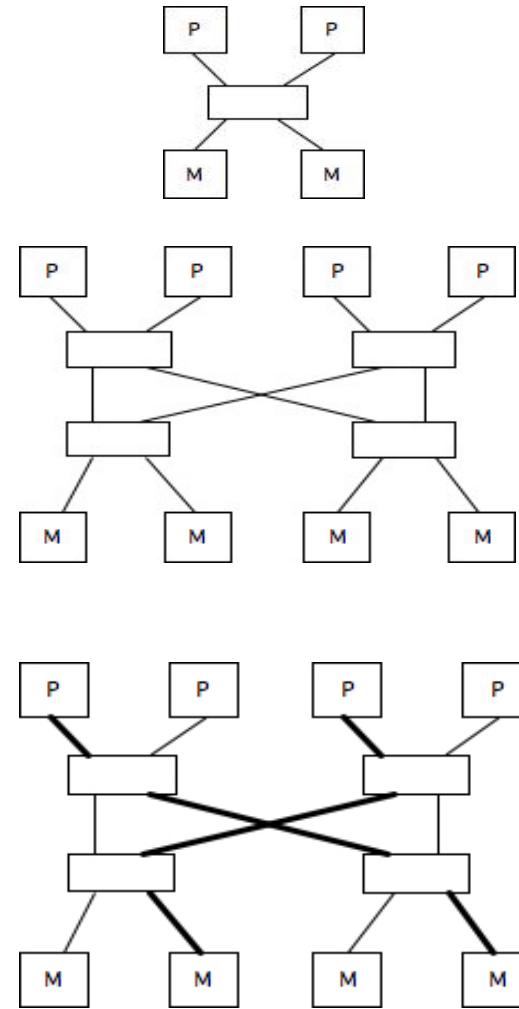


Crossbar Switches: Every processor connects to the switch which routes communications to their destinations



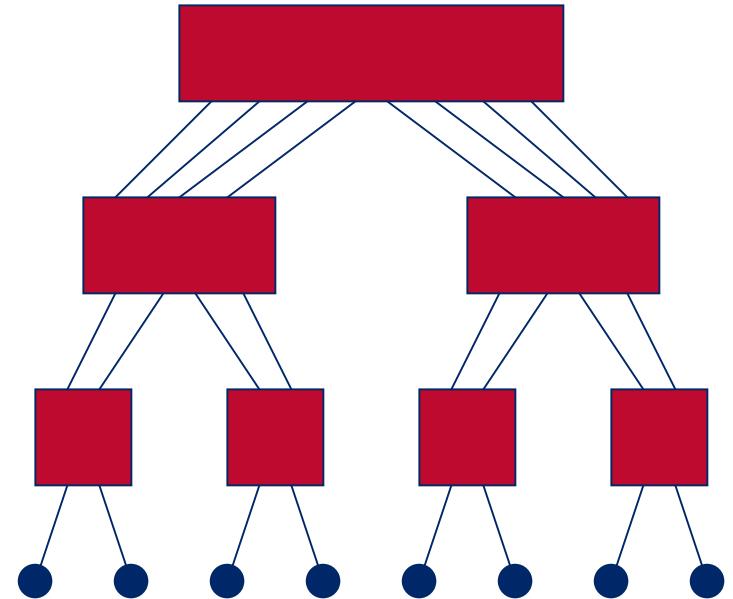
Butterfly exchange network

- Built out of simple switching elements
- Multi-stage; #stages grows with #procs
- Multiple non-colliding paths possible
- Uniform memory access



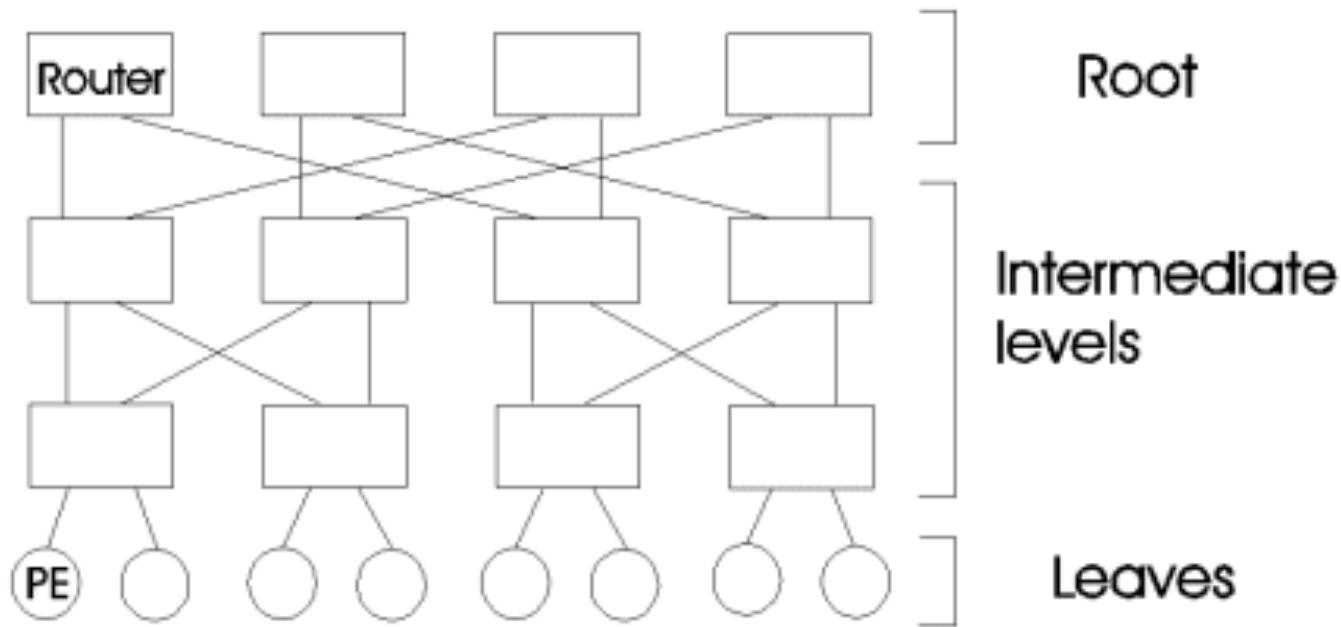
Fat Trees

- Multiple switches
- Each level has the same number of links in as out
- Increasing number of links at each level
- Gives full bandwidth between the links
- Added latency the higher you go



Fat Trees

- in practice emulated by switching network



Interconnect graph theory

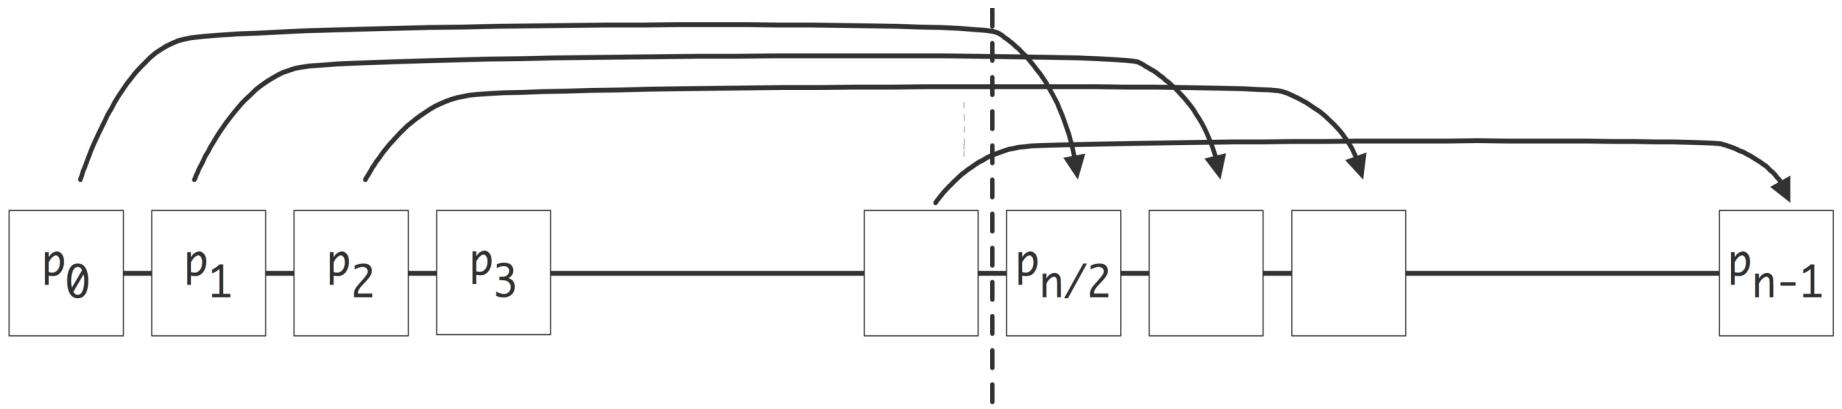
- Degree
 - How many links to other processors does each node have?
 - More is better, but also expensive and hard to engineer
- Diameter
 - maximum distance between any two processors in the network.
 - The distance between two processors is defined as the shortest path, in terms of links, between them.
 - completely connected network is 1, for star network is 2, for ring is $p/2$ (for p even processors)
- Connectivity
 - measure of the multiplicity of paths between any two processors (# arcs that must be removed to break the connection).
 - high connectivity is desired since it lowers contention for communication resources.
 - 1 for linear array, 1 for star, 2 for ring, 2 for mesh, 4 for torus
 - technically 1 for traditional fat trees, but there is redundancy in the switch infrastructure

Practical issues in interconnects

- Latency : How long does it take to start sending a "message"? Units are generally microseconds or milliseconds.
- Bandwidth : What data rate can be sustained once the message is started? Units are Mbytes/sec or Gbytes/sec.
 - Both point-to-point and aggregate bandwidth are of interest
- Multiple wires: multiple latencies, same bandwidth
- Sometimes shortcuts possible: `wormhole routing'

Measures of bandwidth

- Aggregate bandwidth: total data rate if every processor sending: total capacity of the wires. This can be very high and quite unrealistic.
- Imagine linear array with processor i sending to $P/2+i$: ‘Contention’
- Bisection bandwidth: bandwidth across the minimum number of wires that would split the machine in two.



Interconnects

- **Bisection width**
 - Minimum # of communication links that have to be removed to partition the network into two equal halves. Bisection width is
 - 2 for ring, $\text{sq. root}(p)$ for mesh with p (even) processors, $p/2$ for hypercube, $(p*p)/4$ for completely connected (p even).
- **Channel width**
 - of physical wires in each communication link
- **Channel rate**
 - peak rate at which a single physical wire link can deliver bits
- **Channel BW**
 - peak rate at which data can be communicated between the ends of a communication link
 - = (channel width) * (channel rate)
- **Bisection BW**
 - minimum volume of communication found between any 2 halves of the network with equal # of procs
 - = (bisection width) * (channel BW)

Parallel programming



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

Programming the memory models

- Shared memory: all processors share the same address space
 - OpenMP: directives-based programming
 - PGAS languages (UPC, Titanium, X10)
- Distributed memory: every processor has its own address space
 - MPI: Message Passing Interface

Ideal vs Practice

- Shared memory (or SMP: Symmetric MultiProcessor) is easy to program (OpenMP) but hard to build
 - bus-based systems can become saturated
 - large, fast (high bandwidth, low latency) crossbars are expensive
 - cache-coherency is hard to maintain at scale

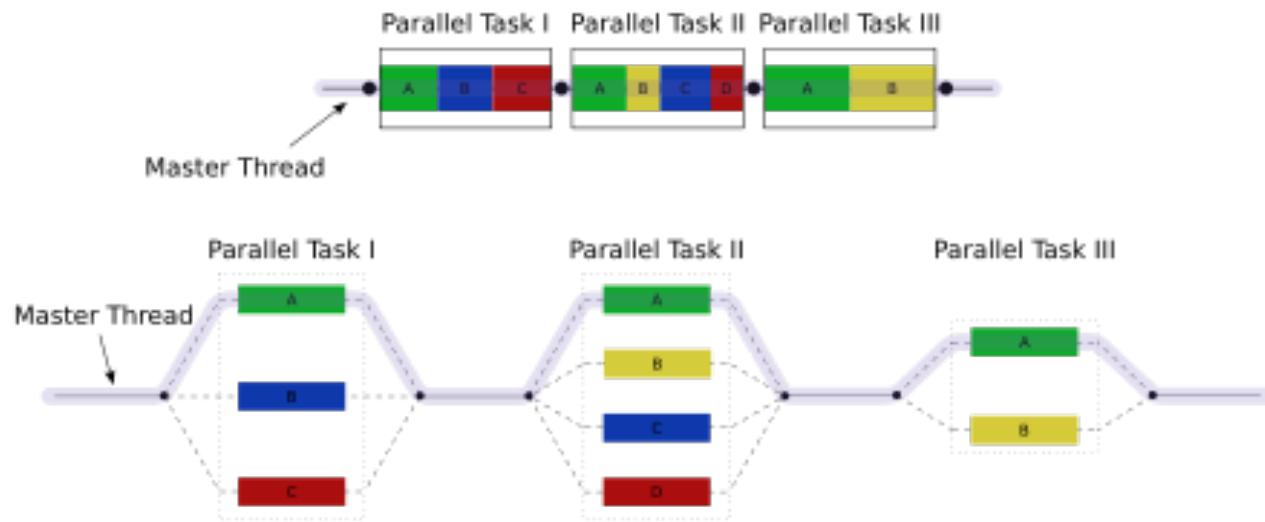
Ideal vs Practice

- Distributed memory is easy to build (bunch of PCs, ethernet) but hard to program (MPI)
 - You have to spell it all out
 - interconnects have higher latency, so data is not immediately there
 - makes parallel algorithm development and programming harder

Programmer's view vs Hard reality

- It is possible for distributed hardware to act like shared
- Middle layer: programmatic, OS, hardware support
- New machines: SGI UV, Cray Gemini

Shared memory programming in OpenMP



- Shared memory.
- Various issues: critical regions, binding, thread overhead

Thread programming

- Threads have shared address space (unlike processes)
- Great for parallel processing on shared memory
- Ex: quad-core => use 4 threads (8 with HT)
- OpenMP declares parallel tasks, the threads execute them in some order (shared memory essential!)
- Obvious example: loop iterations can be parallel

OpenMP programming

- “pragma”-based: directives to the compiler

```
#pragma omp parallel default(none) \
    shared(n,x,y) private(i)
{
    #pragma omp for
    for (i=0; i<n; i++)
        x[i] += y[i];
} /*-- End of parallel region --*/
```

clauses

```
!$omp parallel default(none) &
 !$omp shared(n,x,y) private(i)
 !$omp do
     do i = 1, n
         x(i) = x(i) + y(i)
     end do
 !$omp end do
 !$omp end parallel
```

OpenMP programming

- Handling of private and shared data

```
sum = 0.0
!$omp parallel default(none) &
!$omp shared(n,x) private(i)
!$omp do reduction (+:sum)
do i = 1, n
    sum = sum + x(i)
end do
!$omp end do
!$omp end parallel
print *,sum
```

Variable *SUM* is a
shared variable

Now that threads have come up...

- Your typical core can handle one thread (two with HT)
- ‘Context switching’ is expensive
- GPU handles many threads with ease, in fact relies on it
- => GPU is even more SIMD than you already realized

On to Distributed Memory



THE UNIVERSITY OF TEXAS AT AUSTIN
TEXAS ADVANCED COMPUTING CENTER

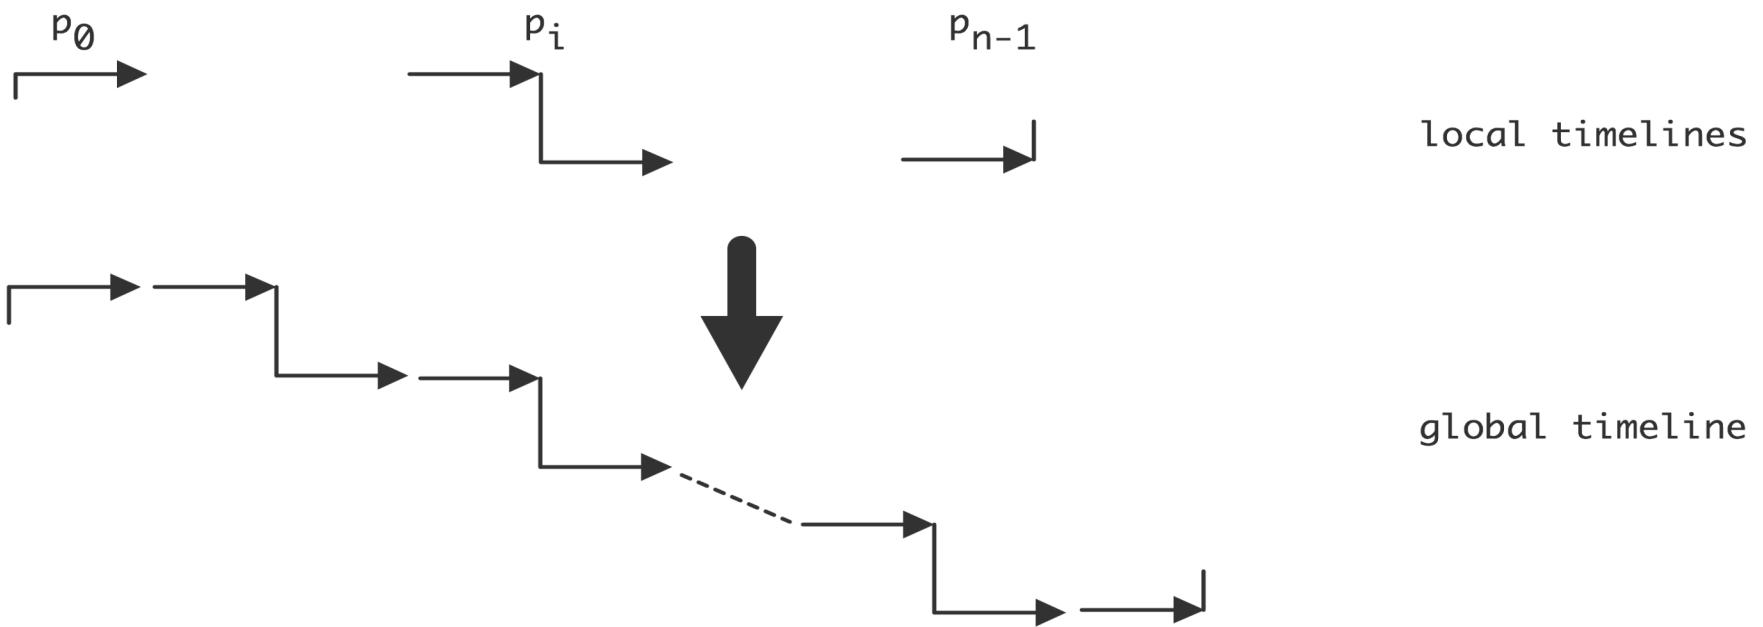
Parallel algorithms vs parallel programming

- Example: two arrays x and y ; n processors; p_i stores x_i and y_i
- Algorithm: $y_i := y_i + x_{i-1}$
- Global description:
 - Processors $0..n-2$ send their x element to the right
 - Processors $1..n-1$ receive an x element from the left
 - Add the received number to their y element

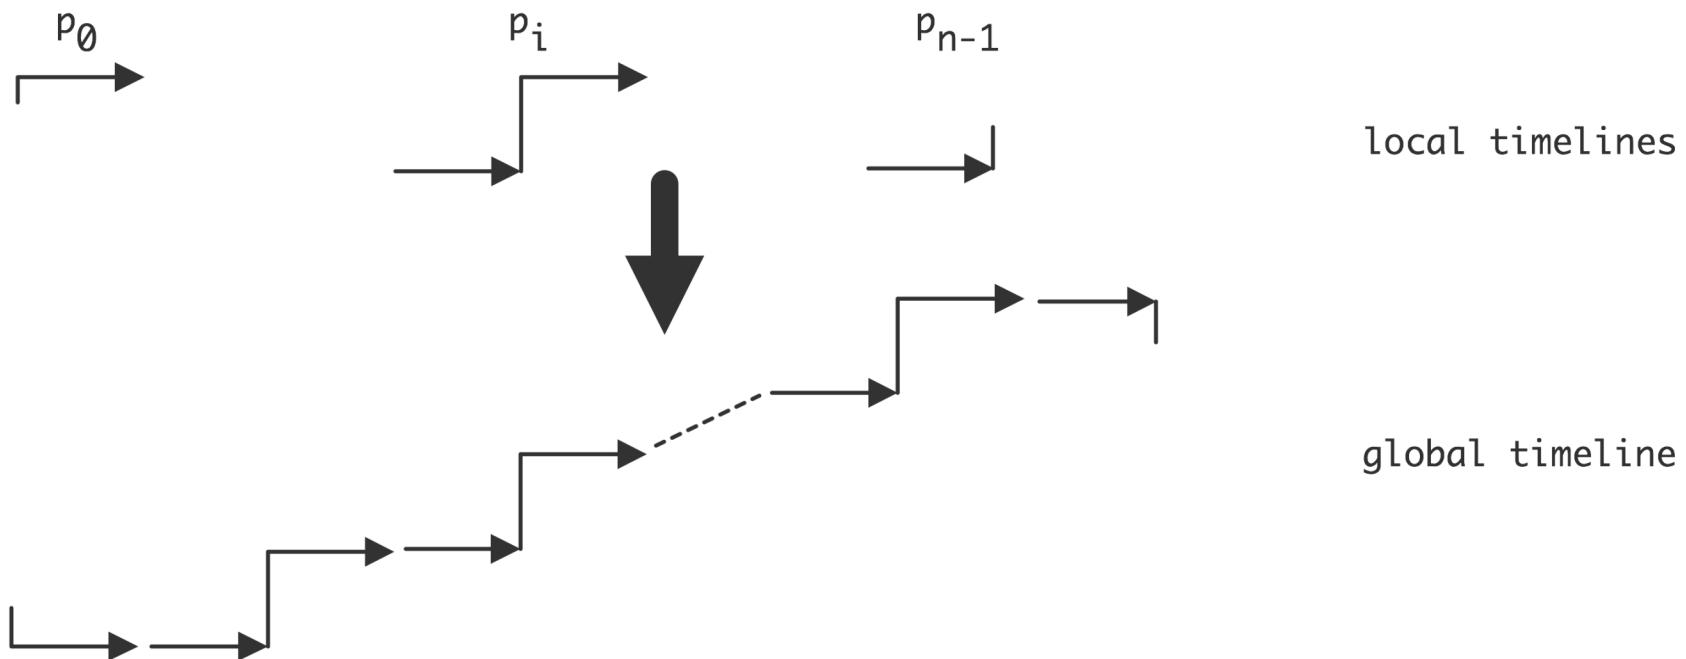
Local implementations

- One implementation:
 - If my number >0 : receive a x element, add it to my y element
 - If my number $<n-1$: send my x element to the right
- Other implementation
 - If my number $<n-1$: send my x element to the right
 - If my number >0 : receive a x element, add it to my x element

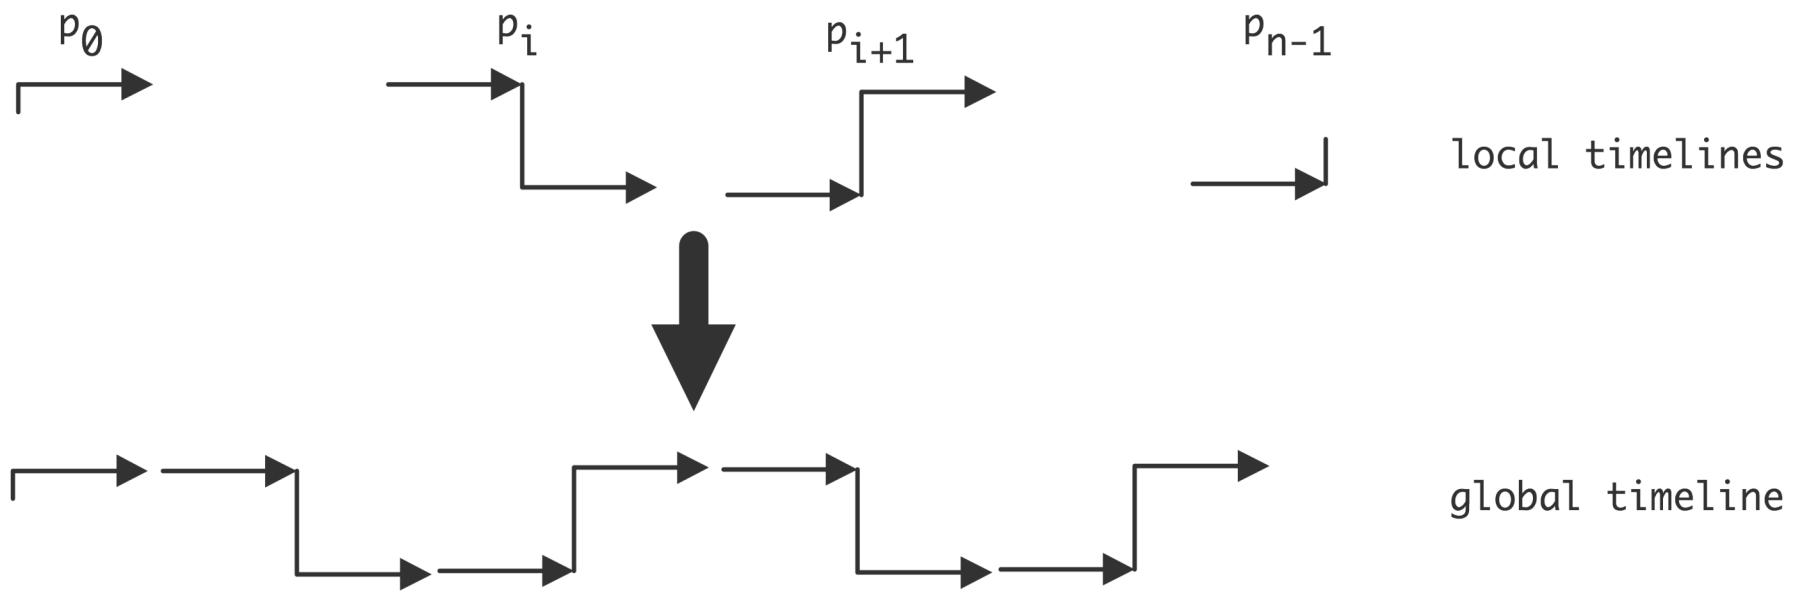
- One implementation:
 - If my number >0 : receive a x element, add it to my y element
 - If my number $<n-1$: send my x element to the right



- Other implementation
 - If my number $< n-1$: send my x element to the right
 - If my number > 0 : receive a x element, add it to my y element



- Better implementation
 - If my number odd: receive then send
 - If my number even: send then receive



Blocking operations

- Send & recv operations are *blocking*: a send does not finished until the message is actually received
- Parallel operation becomes sequentialized; in a ring even loads to *deadlock*

Non-Blocking operations

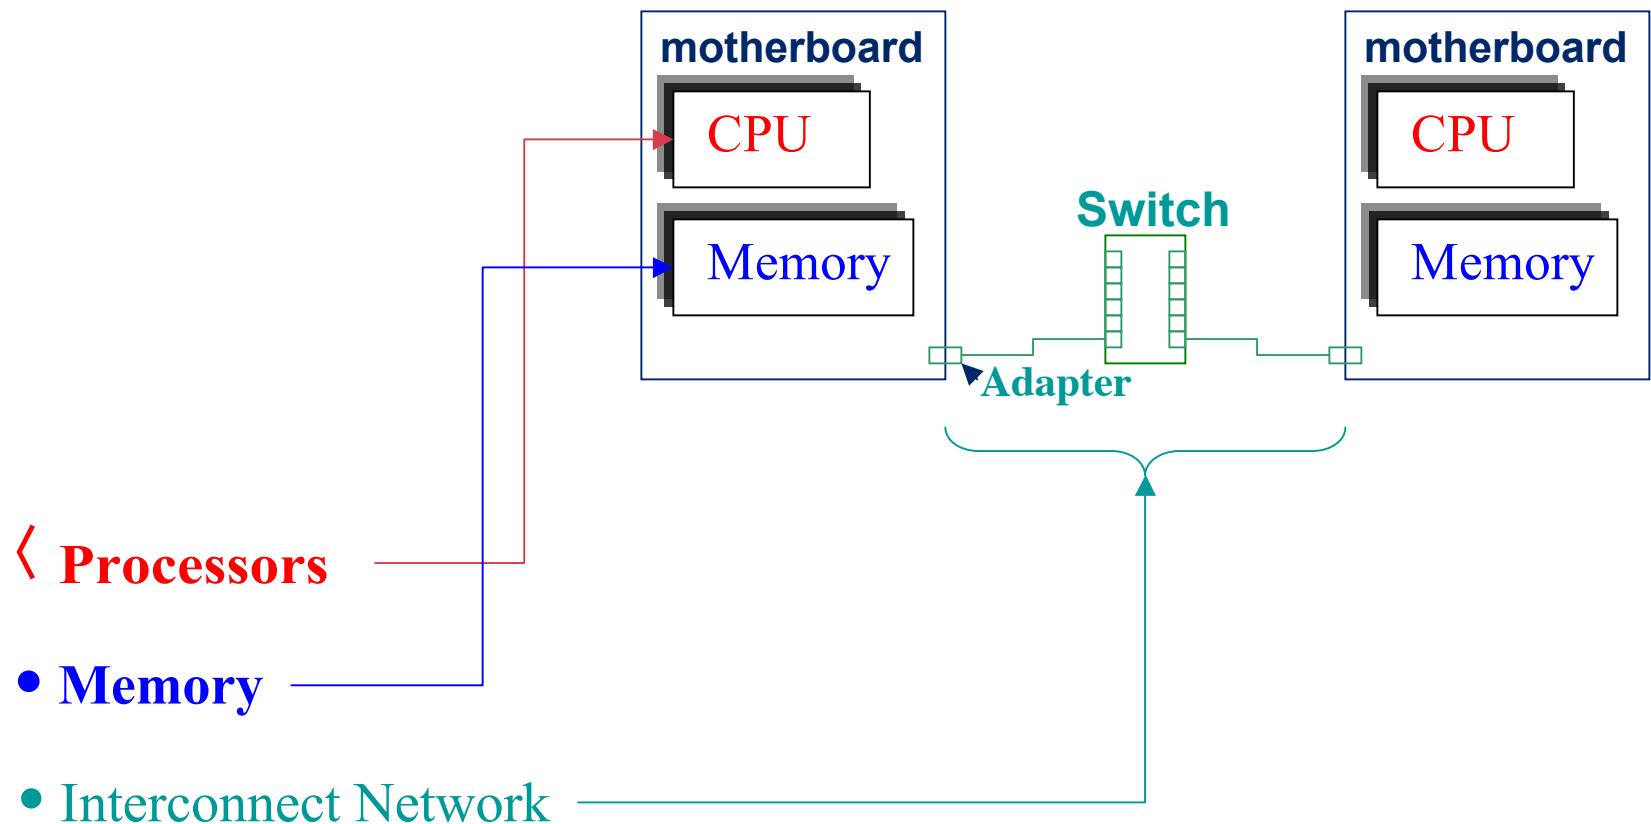
- Non-blocking send & recv:
 - Give a buffer to the system to send from / recv into
 - Continue with next instruction
 - Check for completion later

MPI: message passing

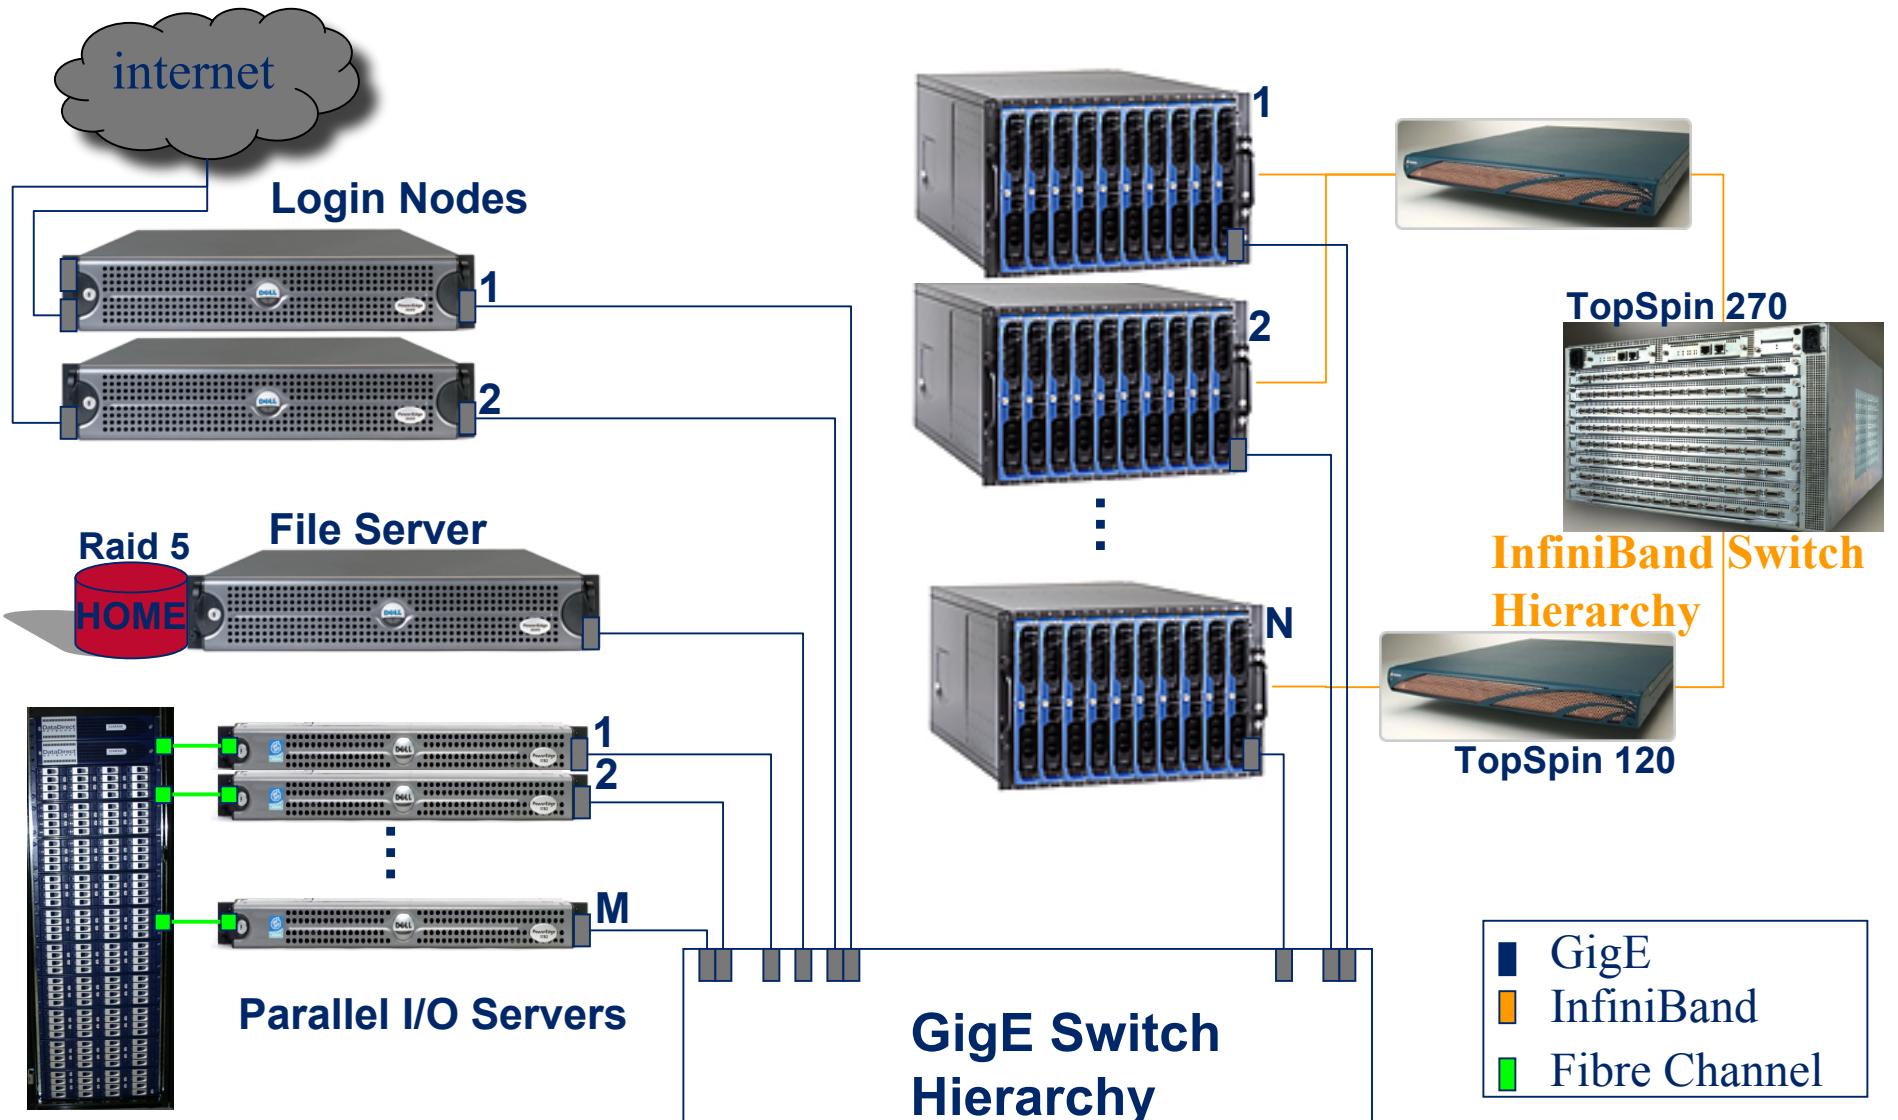
- Message Passing Interface: library for explicit communication
- Point-to-point and collective communication
- Blocking semantics, buffering
- Looks harder than it is

```
if(myid == 0)
{
    printf("WE have %d processors\n", numprocs);
    for(i=1;i<numprocs;i++)
    {
        sprintf(buff, "Hello %d", i);
        MPI_Send(buff, 128, MPI_CHAR,
                 i, 0, MPI_COMM_WORLD);
    }
    for(i=1;i<numprocs;i++)
    {
        MPI_Recv(buff, 128, MPI_CHAR,
                 i, 0, MPI_COMM_WORLD, &stat);
        printf("%s\n", buff);
    }
}
else
{
    MPI_Recv(buff, 128, MPI_CHAR,
             0, 0, MPI_COMM_WORLD, &stat);
    sprintf(idstr, " Processor %d ", myid);
    strcat(buff, idstr);
    strcat(buff, "reporting for duty\n");
    MPI_Send(buff, 128, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
}
```

Basic Anatomy of a Server/Desktop/Laptop/Cluster-node node



Lonestar @ TACC



RAID

- Was: Redundant Array of Inexpensive Disks
- Now: Redundant Array of Independent Disks
- Multiple disk drives working together to:
 - increase capacity of a single logical volume
 - increase performance
 - improve reliability/add fault tolerance
- 1 Server with RAIDed disks can provide disk access to multiple nodes with NFS

Parallel Filesystems

- Use multiple servers together to aggregate disks
 - utilizes RAIDed disks
 - improved performance
 - even higher capacities
 - may use high-performance network
- Vendors/Products
 - CFS/Lustre
 - IBM/GPFS
 - IBRIX/IBRIXFusion
 - RedHat/GFS
 - ...

Summary

- Why so much parallel talk?
 - Every computer is a parallel computer now
 - Good serial computing skills a central to good parallel computing
 - Cluster and MPP nodes appear largely like desktops and laptops
 - Processing units: CPUs, FPUs, GPUs
 - Memory hierarchies: Registers, Caches, Main memory
 - Internal Interconnect: Buses and Switch-based networks
 - Clusters and MPPs built via fancy connections.